

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 01:27 | Laufzeit: 2h 22m | Run: model_comparison_final

Metrik	LSTM	Transformer	MLP
Steps	250,000	250,000	250,000
Fortschritt	0.83%	0.83%	0.83%
Reward (aktuell)	-20.445	-20.005	-19.176
Curriculum Phase	1	1	1
Episode in Phase	2,125	2,122	2,120
Gate-SuccessRate	5.0%	4.9%	5.0%
Rolling Success	5.6%	2.8%	2.4%
Beste Success	48.0%	32.3%	37.7%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.000	0.000	0.000
Sprünge/Episode	31.829	29.524	30.026
Tod: Lava	0.0%	0.0%	0.0%
Tod: Timeout	97.6%	100.0%	97.4%

Reward & Erfolgsrate

Kumulativer Reward



Rolling Success Rate (%)



Curriculum & Todesursachen


Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?

PBR-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken

Policy Loss



Value Loss



Entropy



Learning Rate

