


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 02:27 | Laufzeit: 3h 22m | Run: model_comparison_final


Metrik	LSTM	Transformer	MLP
Steps	640,000	640,000	640,000
Fortschritt	2.13%	2.13%	2.13%
Reward (aktuell)	3.785	-7.131	2.735
Curriculum Phase	1	1	1
Episode in Phase	632	601	642
Gate-SuccessRate	12.9%	13.0%	12.9%
Rolling Success	26.8%	3.5%	20.2%
Beste Success	48.0%	32.3%	37.7%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.416	0.338	0.388
Sprünge/Episode	5.199	14.225	9.127
Tod: Lava	78.3%	57.7%	73.9%
Tod: Timeout	5.0%	39.4%	9.7%

Reward & Erfolgsrate

Kumulativer Reward



Rolling Success Rate (%)




Curriculum & Todesursachen


Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?

PBRS-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

