


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 04:28 | Laufzeit: 5h 22m | Run: model_comparison_final


Metrik	LSTM	Transformer	MLP
Steps	1,350,000	1,350,000	1,350,000
Fortschritt	4.50%	4.50%	4.50%
Reward (aktuell)	3.046	-2.832	8.660
Curriculum Phase	4	4	4
Episode in Phase	6,612	6,608	6,602
Gate-SuccessRate	18.8%	18.8%	18.7%
Rolling Success	17.7%	2.7%	24.4%
Beste Success	48.0%	32.3%	37.7%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.368	0.163	0.524
Sprünge/Episode	14.299	9.977	9.943
Tod: Lava	39.1%	40.7%	25.7%
Tod: Timeout	5.7%	7.0%	1.0%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)




Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?

PBRS-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

