


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 05:28 | Laufzeit: 6h 23m | Run: model_comparison_final

Metrik	LSTM	Transformer	MLP
Steps	1,710,000	1,710,000	1,710,000
Fortschritt	5.70%	5.70%	5.70%
Reward (aktuell)	2.498	-3.126	2.892
Curriculum Phase	5	5	5
Episode in Phase	3,577	3,572	3,563
Gate-SuccessRate	12.1%	12.0%	12.0%
Rolling Success	14.9%	1.6%	20.8%
Beste Success	48.0%	32.3%	37.7%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.500	0.172	0.468
Sprünge/Episode	12.670	7.379	8.431
Tod: Lava	38.6%	46.0%	37.6%
Tod: Timeout	2.3%	3.4%	0.0%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)




Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRS-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

