


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 06:28 | Laufzeit: 7h 24m | Run: model_comparison_final


Metrik	LSTM	Transformer	MLP
Steps	2,050,000	2,050,000	2,050,000
Fortschritt	6.83%	6.83%	6.83%
Reward (aktuell)	3.840	-3.305	8.379
Curriculum Phase	5	5	5
Episode in Phase	8,551	8,540	8,534
Gate-SuccessRate	16.3%	16.0%	15.9%
Rolling Success	16.6%	1.5%	27.4%
Beste Success	48.0%	32.3%	37.7%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.440	0.165	0.447
Sprünge/Episode	9.670	5.612	9.155
Tod: Lava	35.8%	50.6%	32.0%
Tod: Timeout	0.9%	2.4%	1.0%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)




Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

