


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 10:54 | Laufzeit: 11h 50m | Run: model_comparison_final


Metrik	LSTM	Transformer	MLP
Steps	3,000,000	3,010,000	3,010,000
Fortschritt	10.00%	10.03%	10.03%
Reward (aktuell)	-0.284	-4.762	-0.155
Curriculum Phase	6	6	6
Episode in Phase	12,788	12,926	12,926
Gate-SuccessRate	8.0%	7.2%	7.2%
Rolling Success	7.5%	0.8%	9.1%
Beste Success	48.0%	32.3%	37.7%
Lava-Crossings (Summe)	0.017	0.000	0.000
Lava-Sprungversuche	0.635	0.181	0.657
Sprünge/Episode	10.654	8.417	11.273
Tod: Lava	39.4%	34.7%	50.5%
Tod: Timeout	0.0%	4.2%	0.0%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken

Policy Loss



Value Loss



Entropy



Learning Rate

