


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 11:55 | Laufzeit: 12h 49m | Run: model_comparison_final


Metrik	LSTM	Transformer	MLP
Steps	3,370,000	3,370,000	3,370,000
Fortschritt	11.23%	11.23%	11.23%
Reward (aktuell)	0.184	-3.706	0.557
Curriculum Phase	6	6	6
Episode in Phase	18,224	18,203	18,211
Gate-SuccessRate	7.7%	7.6%	7.7%
Rolling Success	8.0%	1.9%	12.4%
Beste Success	48.0%	32.3%	37.7%
Lava-Crossings (Summe)	0.017	0.000	0.000
Lava-Sprungversuche	0.763	0.153	0.514
Sprünge/Episode	15.263	7.141	9.640
Tod: Lava	38.2%	36.5%	40.5%
Tod: Timeout	1.3%	1.2%	0.9%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)




Curriculum & Todesursachen


Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

