


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 12:55 | Laufzeit: 13h 50m | Run: model_comparison_final


Metrik	LSTM	Transformer	MLP
Steps	3,740,000	3,740,000	3,740,000
Fortschritt	12.47%	12.47%	12.47%
Reward (aktuell)	11.265	-1.752	11.766
Curriculum Phase	4	4	4
Episode in Phase	4,135	4,135	4,125
Gate-SuccessRate	32.8%	32.8%	32.6%
Rolling Success	41.6%	6.4%	44.8%
Beste Success	48.0%	32.3%	45.0%
Lava-Crossings (Summe)	0.017	0.000	0.000
Lava-Sprungversuche	0.339	0.200	0.220
Sprünge/Episode	8.009	7.753	6.106
Tod: Lava	20.0%	40.0%	7.3%
Tod: Timeout	2.6%	7.1%	1.6%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)




Curriculum & Todesursachen


Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]




Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

