

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 15-39 | Laufzeit: 0h 58m | Run: model_comparison_final_v2

Metrik	LSTM	Transformer	MLP
Steps	350,000	350,000	350,000
Fortschritt	1.17%	1.17%	1.17%
Reward (aktuell)	28.330	-16.853	28.847
Curriculum Phase	1	1	1
Episode in Phase	1,405	1,434	1,405
Gate-SuccessRate	82.5%	82.9%	82.5%
Rolling Success	94.2%	7.8%	87.7%
Beste Success	94.2%	34.6%	87.7%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.000	0.000	0.000
Sprünge/Episode	6.671	28.581	7.674
Tod: Lava	0.0%	0.0%	0.0%
Tod: Timeout	0.0%	90.7%	0.7%

Reward & Erfolgsrate

Kumulativer Reward



Rolling Success Rate (%)




Curriculum & Todesursachen


Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRS-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken

Policy Loss



Value Loss



Entropy



Learning Rate

