



KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 17-13 | Laufzeit: 2h 31m | Run: model_comparison_final_v2

Metrik	LSTM	Transformer	MLP
Steps	690,000	690,000	690,000
Fortschritt	2.30%	2.30%	2.30%
Reward (aktuell)	14.490	-6.857	14.836
Curriculum Phase	4	4	4
Episode in Phase	1,766	1,770	1,741
Gate-SuccessRate	46.7%	46.4%	47.2%
Rolling Success	58.9%	3.2%	51.6%
Beste Success	94.2%	34.6%	91.8%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.395	0.408	0.381
Sprünge/Episode	6.790	14.132	6.714
Tod: Lava	31.1%	47.4%	35.1%
Tod: Timeout	4.8%	19.7%	5.4%

Reward & Erfolgsrate



Curriculum & Todesursachen


Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)



Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

