

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 19-13 | Laufzeit: 4h 32m | Run: model_comparison_final_v2


Metrik	LSTM	Transformer	MLP
Steps	1,440,000	1,440,000	1,440,000
Fortschritt	4.80%	4.80%	4.80%
Reward (aktuell)	13.762	1.572	11.766
Curriculum Phase	5	5	5
Episode in Phase	5,747	5,772	5,769
Gate-SuccessRate	42.1%	42.2%	42.2%
Rolling Success	52.5%	21.7%	45.8%
Beste Success	94.2%	34.6%	91.8%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.453	0.495	0.422
Sprünge/Episode	10.162	10.611	10.431
Tod: Lava	23.1%	51.6%	30.3%
Tod: Timeout	5.1%	9.5%	11.0%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen

Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRS-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

