

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-09 21:14 | Laufzeit: 6h 33m | Run: model_comparison_final_v2

Metrik	LSTM	Transformer	MLP
Steps	2,190,000	2,190,000	2,190,000
Fortschritt	7.30%	7.30%	7.30%
Reward (aktuell)	12.835	5.123	8.372
Curriculum Phase	6	6	6
Episode in Phase	4,205	4,222	4,171
Gate-SuccessRate	38.9%	39.1%	38.7%
Rolling Success	47.9%	28.2%	36.8%
Beste Success	94.2%	34.6%	91.8%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.476	0.707	0.663
Sprünge/Episode	11.806	11.293	10.730
Tod: Lava	29.1%	50.5%	32.6%
Tod: Timeout	1.9%	9.1%	6.7%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen

Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?


PBRS-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

