

# KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 01:16 | Laufzeit: 10h 34m | Run: model\_comparison\_final\_v2

| Metrik                 | LSTM      | Transformer | MLP       |
|------------------------|-----------|-------------|-----------|
| Steps                  | 3,530,000 | 3,530,000   | 3,530,000 |
| Fortschritt            | 11.77%    | 11.77%      | 11.77%    |
| Reward (aktuell)       | 10.694    | -2.732      | 3.817     |
| Curriculum Phase       | 7         | 7           | 7         |
| Episode in Phase       | 16,898    | 16,889      | 16,922    |
| Gate-SuccessRate       | 34.5%     | 34.5%       | 34.1%     |
| Rolling Success        | 42.6%     | 20.1%       | 26.5%     |
| Beste Success          | 94.2%     | 34.6%       | 91.8%     |
| Lava-Crossings (Summe) | 0.000     | 0.000       | 0.000     |
| Lava-Sprungversuche    | 0.704     | 0.483       | 0.621     |
| Sprünge/Episode        | 15.049    | 16.241      | 15.985    |
| Tod: Lava              | 23.5%     | 43.1%       | 33.3%     |
| Tod: Timeout           | 3.7%      | 17.2%       | 10.6%     |

# Reward & Erfolgsrate

## Kumulativer Reward




## Rolling Success Rate (%)



# Curriculum & Todesursachen


## Curriculum Phase




## Gate-SuccessRate [%] (geteilt)



## Tod durch Timeout [%]




## Tod durch Lava [%]



# Lava-Fokus: Lernt der Agent den Sprung?

## Lava-Sprungversuche / Episode



## Erfolgreiche Überquerungen / Episode



## Sprünge gesamt / Episode



## Distanz zum Ziel bei Episodenende




# Reward-Zerlegung: Woher kommt der Reward?

PBRs-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




# Policy-Metriken

## Policy Loss



## Value Loss



## Entropy



## Learning Rate

