

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 14-31 | Laufzeit: 23h 50m | Run: model_comparison_final_v2

Metrik	LSTM	Transformer	MLP
Steps	7,500,000	7,500,000	7,500,000
Fortschritt	25.00%	25.00%	25.00%
Reward (aktuell)	22.823	7.168	9.985
Curriculum Phase	6	6	6
Episode in Phase	7,988	7,941	8,032
Gate-SuccessRate	58.6%	58.7%	58.9%
Rolling Success	76.2%	40.3%	56.8%
Beste Success	94.2%	52.6%	91.8%
Lava-Crossings (Summe)	0.000	0.066	0.000
Lava-Sprungversuche	0.315	0.443	0.449
Sprünge/Episode	6.165	13.568	11.787
Tod: Lava	11.0%	31.8%	25.8%
Tod: Timeout	2.4%	15.9%	13.5%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?


PBRS-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken

Policy Loss



Value Loss



Entropy



Learning Rate

