


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 15:32 | Laufzeit: 24h 52m | Run: model_comparison_final_v2

Metrik	LSTM	Transformer	MLP
Steps	7,890,000	7,900,000	7,900,000
Fortschritt	26.30%	26.33%	26.33%
Reward (aktuell)	20.807	4.148	10.828
Curriculum Phase	6	6	6
Episode in Phase	6,956	6,991	7,095
Gate-SuccessRate	50.2%	51.5%	51.9%
Rolling Success	67.7%	33.1%	44.2%
Beste Success	94.2%	52.6%	91.8%
Lava-Crossings (Summe)	0.000	0.066	0.000
Lava-Sprungversuche	0.437	0.516	0.500
Sprünge/Episode	8.000	12.505	13.744
Tod: Lava	16.5%	34.7%	20.5%
Tod: Timeout	1.9%	11.6%	11.5%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]




Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?

PBRS-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

