

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-11 04:50 | Laufzeit: 38h 10m | Run: model_comparison_final_v2

Metrik	LSTM	Transformer	MLP
Steps	12,440,000	12,440,000	12,440,000
Fortschritt	41.47%	41.47%	41.47%
Reward (aktuell)	20.340	6.228	8.036
Curriculum Phase	6	6	6
Episode in Phase	9,385	9,477	9,156
Gate-SuccessRate	51.5%	51.3%	52.0%
Rolling Success	72.2%	35.9%	48.1%
Beste Success	94.2%	52.6%	91.8%
Lava-Crossings (Summe)	0.000	0.066	0.000
Lava-Sprungversuche	0.432	0.640	0.533
Sprünge/Episode	8.758	11.780	14.213
Tod: Lava	13.7%	37.0%	18.7%
Tod: Timeout	3.2%	9.0%	13.3%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]




Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?

PBRS-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

