

# KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-11 07:56 | Laufzeit: 41h 16m | Run: model\_comparison\_final\_v2

| Metrik                 | LSTM       | Transformer | MLP        |
|------------------------|------------|-------------|------------|
| Steps                  | 13,490,000 | 13,490,000  | 13,490,000 |
| Fortschritt            | 44.97%     | 44.97%      | 44.97%     |
| Reward (aktuell)       | 22.437     | 6.975       | 9.304      |
| Curriculum Phase       | 6          | 6           | 6          |
| Episode in Phase       | 2,916      | 2,926       | 2,936      |
| Gate-SuccessRate       | 52.8%      | 52.9%       | 52.9%      |
| Rolling Success        | 77.0%      | 36.8%       | 52.6%      |
| Beste Success          | 94.2%      | 54.1%       | 91.8%      |
| Lava-Crossings (Summe) | 0.000      | 0.066       | 0.000      |
| Lava-Sprungversuche    | 0.473      | 0.592       | 0.471      |
| Sprünge/Episode        | 7.745      | 11.158      | 15.600     |
| Tod: Lava              | 12.7%      | 47.5%       | 18.6%      |
| Tod: Timeout           | 0.9%       | 4.2%        | 11.4%      |

# Reward & Erfolgsrate

## Kumulativer Reward



## Rolling Success Rate (%)



# Curriculum & Todesursachen

## Curriculum Phase



## Gate-SuccessRate [%] (geteilt)



## Tod durch Timeout [%]



## Tod durch Lava [%]




# Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende



# Reward-Zerlegung: Woher kommt der Reward?

PBRs-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)



# Policy-Metriken

## Policy Loss



## Value Loss



## Entropy



## Learning Rate

