

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-11 16-12 | Laufzeit: 49h 34m | Run: model_comparison_final_v2

Metrik	LSTM	Transformer	MLP
Steps	16,310,000	16,310,000	16,320,000
Fortschritt	54.37%	54.37%	54.40%
Reward (aktuell)	24.718	14.425	15.903
Curriculum Phase	5	5	5
Episode in Phase	7,627	7,643	4,266
Gate-SuccessRate	66.7%	66.7%	65.8%
Rolling Success	84.6%	42.3%	60.2%
Beste Success	94.2%	54.1%	91.8%
Lava-Crossings (Summe)	0.000	0.066	0.000
Lava-Sprungversuche	0.465	0.543	0.611
Sprünge/Episode	7.561	8.600	9.973
Tod: Lava	6.1%	30.0%	21.2%
Tod: Timeout	1.8%	3.6%	4.4%

Reward & Erfolgsrate

Kumulativer Reward



Rolling Success Rate (%)




Curriculum & Todesursachen


Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]




Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?

PBRS-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken

Policy Loss



Value Loss



Entropy



Learning Rate

