

# KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 00-32 | Laufzeit: 0h 00m | Run: model\_comparison\_final\_v3

| Metrik                 | LSTM    | Transformer | MLP     |
|------------------------|---------|-------------|---------|
| Steps                  | 10,000  | 10,000      | 10,000  |
| Fortschritt            | 0.03%   | 0.03%       | 0.03%   |
| Reward (aktuell)       | -18.703 | -17.580     | -16.350 |
| Curriculum Phase       | 1       | 1           | 1       |
| Episode in Phase       | 439     | 451         | 438     |
| Gate-SuccessRate       | 7.2%    | 6.7%        | 7.3%    |
| Rolling Success        | 0.9%    | 1.6%        | 2.4%    |
| Beste Success          | 0.9%    | 1.6%        | 2.4%    |
| Lava-Crossings (Summe) | 0.000   | 0.000       | 0.000   |
| Lava-Sprungversuche    | 0.000   | 0.000       | 0.000   |
| Sprünge/Episode        | 16.163  | 17.118      | 15.880  |
| Tod: Lava              | 0.0%    | 0.0%        | 0.0%    |
| Tod: Timeout           | 47.8%   | 51.0%       | 47.8%   |

# Reward & Erfolgsrate

Kumulativer Reward



Rolling Success Rate (%)



# Curriculum & Todesursachen

Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



# Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende



# Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)



