


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 01:32 | Laufzeit: 1h 00m | Run: model_comparison_final_v3

Metrik	LSTM	Transformer	MLP
Steps	1,110,000	1,110,000	1,110,000
Fortschritt	3.70%	3.70%	3.70%
Reward (aktuell)	11.743	-5.566	19.730
Curriculum Phase	4	4	3
Episode in Phase	2,315	2,297	2,326
Gate-SuccessRate	44.1%	44.3%	44.0%
Rolling Success	34.2%	3.5%	69.8%
Beste Success	56.1%	4.0%	73.4%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.586	0.393	0.341
Sprünge/Episode	8.657	13.933	5.816
Tod: Lava	52.1%	58.4%	20.1%
Tod: Timeout	2.1%	19.1%	6.7%

Reward & Erfolgsrate

Kumulativer Reward



Rolling Success Rate (%)



Curriculum & Todesursachen


Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRS-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)



Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

