


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 03:32 | Laufzeit: 3h 00m | Run: model_comparison_final_v3


Metrik	LSTM	Transformer	MLP
Steps	3,140,000	3,140,000	3,140,000
Fortschritt	10.47%	10.47%	10.47%
Reward (aktuell)	8.682	5.012	18.449
Curriculum Phase	6	6	6
Episode in Phase	4,216	4,540	4,529
Gate-SuccessRate	48.3%	48.3%	48.5%
Rolling Success	36.1%	34.2%	72.4%
Beste Success	56.1%	41.1%	76.3%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.574	0.617	0.153
Sprünge/Episode	11.535	15.628	5.135
Tod: Lava	31.7%	37.2%	12.6%
Tod: Timeout	6.9%	14.9%	3.6%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen


Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]




Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?


Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRS-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)




Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

