


KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 04:32 | Laufzeit: 4h 00m | Run: model_comparison_final_v3


Metrik	LSTM	Transformer	MLP
Steps	4,150,000	4,160,000	4,150,000
Fortschritt	13.83%	13.87%	13.83%
Reward (aktuell)	8.564	4.019	18.339
Curriculum Phase	6	6	6
Episode in Phase	5,795	5,870	5,762
Gate-SuccessRate	36.3%	36.6%	36.4%
Rolling Success	27.4%	25.6%	59.6%
Beste Success	56.1%	41.1%	76.3%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.492	0.582	0.245
Sprünge/Episode	8.615	12.848	6.255
Tod: Lava	35.2%	34.2%	13.3%
Tod: Timeout	0.8%	12.7%	5.1%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen


Curriculum Phase




Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?

PBRs-Shaping-Summe / Episode




Line-of-Sight-Reward-Summe / Episode



Curiosity Reward



Distanz-Fortschritt [%] (Start→Ende)



Policy-Metriken

Policy Loss



Value Loss



Entropy



Learning Rate

