

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 16-34 | Laufzeit: 16h 02m | Run: model_comparison_final_v3

Metrik	LSTM	Transformer	MLP
Steps	16,020,000	16,020,000	16,020,000
Fortschritt	53.40%	53.40%	53.40%
Reward (aktuell)	11.962	13.838	22.573
Curriculum Phase	6	6	6
Episode in Phase	5,619	5,707	5,629
Gate-SuccessRate	61.3%	61.4%	61.2%
Rolling Success	52.0%	54.0%	77.3%
Beste Success	66.9%	67.0%	88.6%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.619	0.441	0.134
Sprünge/Episode	9.876	6.144	7.103
Tod: Lava	33.6%	24.6%	8.2%
Tod: Timeout	0.9%	2.5%	6.2%

Reward & Erfolgsrate

Kumulativer Reward



Rolling Success Rate (%)




Curriculum & Todesursachen


Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende




Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)



Policy-Metriken


Policy Loss



Value Loss



Entropy



Learning Rate

