

KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-10 18-30 | Laufzeit: 17h 58m | Run: model_comparison_final_v3

Metrik	LSTM	Transformer	MLP
Steps	16,420,000	16,420,000	16,420,000
Fortschritt	54.73%	54.73%	54.73%
Reward (aktuell)	14.568	17.063	20.609
Curriculum Phase	6	6	5
Episode in Phase	2,593	2,616	2,514
Gate-SuccessRate	64.5%	64.6%	64.4%
Rolling Success	45.3%	45.9%	63.8%
Beste Success	66.9%	67.0%	88.6%
Lava-Crossings (Summe)	0.000	0.000	0.000
Lava-Sprungversuche	0.713	0.453	0.194
Sprünge/Episode	10.426	5.556	7.376
Tod: Lava	27.8%	23.1%	8.6%
Tod: Timeout	1.7%	0.9%	8.6%

Reward & Erfolgsrate

Kumulativer Reward




Rolling Success Rate (%)



Curriculum & Todesursachen


Curriculum Phase



Gate-SuccessRate [%] (geteilt)



Tod durch Timeout [%]



Tod durch Lava [%]



Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode




Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode




Distanz zum Ziel bei Episodenende



Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)



Policy-Metriken

Policy Loss



Value Loss



Entropy



Learning Rate

