

# KI-Agenten Trainingsanalyse - Modellvergleich (final)

Stand: 2026-07-11 06:48 | Laufzeit: 30h 16m | Run: model\_comparison\_final\_v3

| Metrik                 | LSTM       | Transformer | MLP        |
|------------------------|------------|-------------|------------|
| Steps                  | 29,040,000 | 29,050,000  | 29,040,000 |
| Fortschritt            | 96.80%     | 96.83%      | 96.80%     |
| Reward (aktuell)       | 14.937     | 16.484      | 15.685     |
| Curriculum Phase       | 7          | 7           | 7          |
| Episode in Phase       | 9,708      | 9,714       | 9,640      |
| Gate-SuccessRate       | 57.5%      | 57.5%       | 57.4%      |
| Rolling Success        | 55.9%      | 63.3%       | 64.7%      |
| Beste Success          | 74.3%      | 79.0%       | 88.6%      |
| Lava-Crossings (Summe) | 0.000      | 0.000       | 0.000      |
| Lava-Sprungversuche    | 1.056      | 1.198       | 0.228      |
| Sprünge/Episode        | 16.352     | 7.341       | 10.649     |
| Tod: Lava              | 26.8%      | 20.9%       | 15.8%      |
| Tod: Timeout           | 5.6%       | 1.1%        | 17.5%      |

# Reward & Erfolgsrate

## Kumulativer Reward



## Rolling Success Rate (%)



# Curriculum & Todesursachen


## Curriculum Phase



## Gate-SuccessRate [%] (geteilt)



## Tod durch Timeout [%]



## Tod durch Lava [%]



# Lava-Fokus: Lernt der Agent den Sprung?

Lava-Sprungversuche / Episode



Erfolgreiche Überquerungen / Episode



Sprünge gesamt / Episode



Distanz zum Ziel bei Episodenende




# Reward-Zerlegung: Woher kommt der Reward?


PBRs-Shaping-Summe / Episode



Line-of-Sight-Reward-Summe / Episode



Curiosity Reward




Distanz-Fortschritt [%] (Start→Ende)



# Policy-Metriken

## Policy Loss



## Value Loss



## Entropy



## Learning Rate

